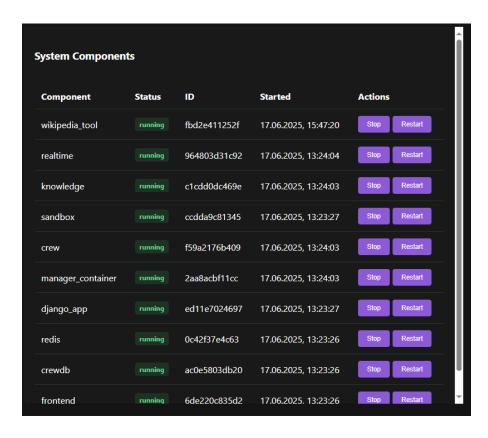
- o Start System
- o Update
- ∘ Stop System





State: Default

State: State: Running Default

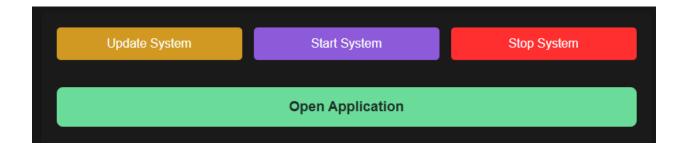
Export Database Import Database

The import happens directly in the browser.



Running

Open Application





Update

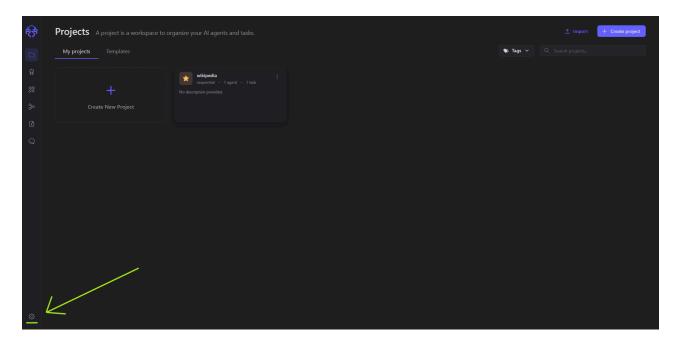
Start System

LLM Settings: How to set up models (Required step)

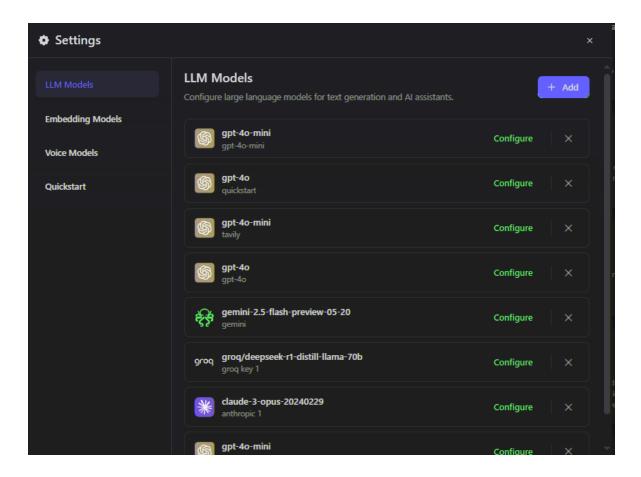
Before you can use agents, flows, tools, or any intelligent features — you **must** set up your LLM models (Large Language Models). This is a required step.

Where to find it

- 1. Go to the **bottom-left corner** of the screen.
- 2. Click the gear icon (Settings).



- 3. A modal window will appear with 4 tabs:
 - LLM Models
 - Embedding Models
 - Voice Models
 - Quickstart



These tabs are where you create and manage all the models needed in your system.

LLM Models tab

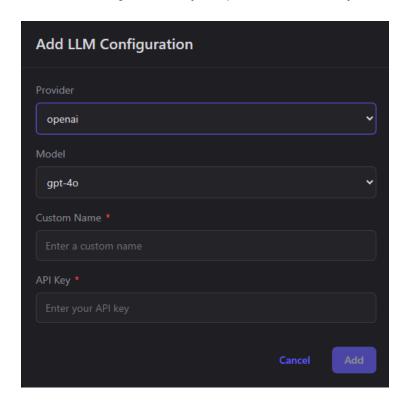
How to Add a New LLM

1. Click the + Add button on the LLM Models tab



- 2. A modal will open with the following required fields:
 - **Provider** Select from a dropdown (e.g., OpenAl, Anthropic, etc.)
 - Model Choose the model name (options depend on the selected provider)
 - Custom Name Give your model a name (you'll use this name later)

API Key – Paste your provider's API key



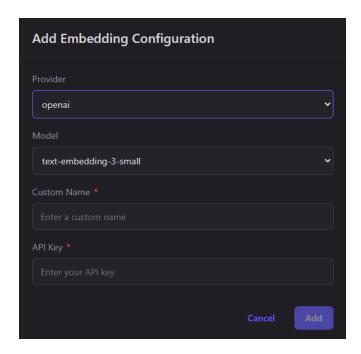
3. After filling out everything, click **Add.** The model will appear in the list of all available LLMs.

Embedding Models tab

This tab works **exactly like LLM Models**, but for **embedding models** (used for memory, semantic search, etc.).

To add one:

- Click + Add
- Fill out: Provider, Model, Custom Name, API Key

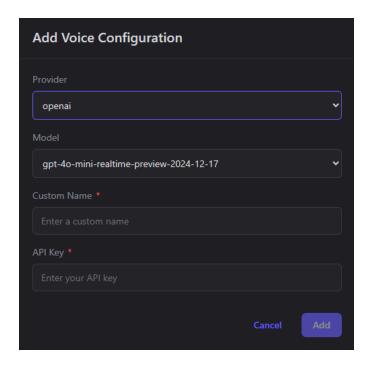


 $\bullet \quad \text{Save} \to \text{The model appears in your list}$

Voice Models tab

Again, this tab works the same way. Voice models are used for **speech synthesis** (e.g. if your agent speaks aloud in a real-time chat).

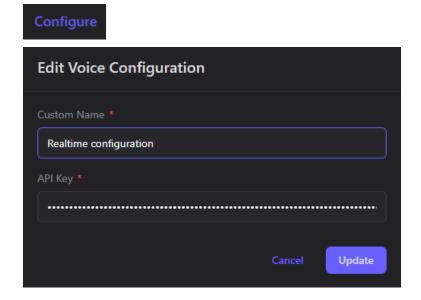
Just click + Add and follow the same steps.



What you can do with each model in the list

Once you've added models, you can:

• Edit / Configure (change name or API key)

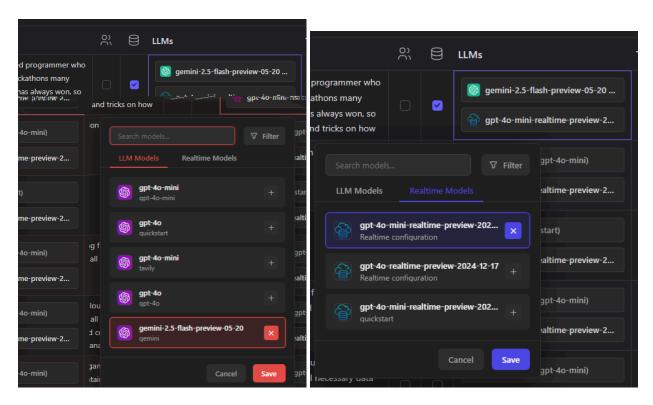


• **Delete** (remove from the system)

How to assign a model to an agent

To make agents actually use a model:

- 1. Go to the **Staff page** (where your agents are listed)
- 2. Find the cell under the column for LLM Model or Realtime Model
- Double-click the cell
 A small window will appear → select the model you want to assign



Quickstart Setup (OpenAl)

The **Quickstart** tab in the Settings panel helps you rapidly configure your environment using a single OpenAl API key.

What It Does

By entering your **OpenAl API Key**, the system will **automatically create and configure** everything you need to start building immediately:

- LLM models and tool configurations
- Realtime model settings
- Embedding model settings
- Default models for:
 - o Projects
 - Agents
 - Tools

This is the **fastest way to get started**, especially useful for first-time users or quick testing.

How to Use It

- 1. Paste your **OpenAl API Key** into the input field.
- 2. Click "Start Building".
- 3. The platform will pre-fill all necessary model and tool settings for you.

⚠ You can still manually adjust or expand these configurations later under the LLM Models, Embedding Models, or Voice Models tabs.

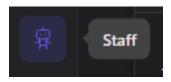
Why This Is Important

Nothing in your system will work without assigning models. Flows will not run, agents will not respond, and real-time chats won't function unless:

- At least one LLM model is added
- It is enabled and properly assigned to agents

Staff in EpicStaff are your workers, agents — virtual helpers, who will bring to life any idea you have. This is where you define the intelligent agents that will perform tasks in your project.

When you open the Agents page, you'll see a table listing all the agents in your project. Each **row** represents a separate **agent**. Each **column** contains a specific property of the agent, which you can edit directly in the table by **doubleclicking** the cell you want to modify.



There are two ways to create and manage agents:

- 1. Using the **Table View**, which allows quick editing of multiple agents in a spreadsheet-like format.
- 2. Using the **Form View**, which provides a step-by-step interface for creating or editing one agent at a time.

1. Creating Agents via the Table View format

Below is a breakdown of each column and its purpose:

Agent Role

A short name or title that describes the function of the agent (e.g., "English Teacher", "Trip Planner", "John Swift").

Goal

A text describing what this agent is supposed to accomplish.

"Search for the cheapest flights that match the user's preferences."

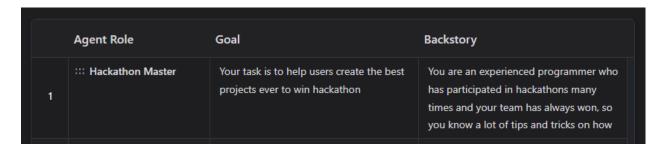
This helps the system understand the agent's purpose and align tasks accordingly.

Backstory

A narrative-style background for the agent.

This helps the system build more human-like, context-aware reasoning by imagining the agent as a "character."

"A former travel agent who now uses AI to help people plan trips faster and smarter."

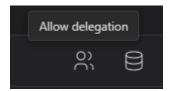


Allow Delegation

A checkbox that enables the agent to delegate tasks to other agents.

If checked, the agent can decide to pass on subtasks to other agents in the system or call other agents for help while performing its task.

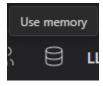
Useful for agents responsible for complex tasks (e.g., "Trip Planner" can delegate hotel search to another agent).



Memory

A checkbox that enables the agent to store all information it gathers or generates during a single flow run. This memory is accessible to the user after execution and can be reviewed or deleted as needed.

To find out more about memory go here: Memory

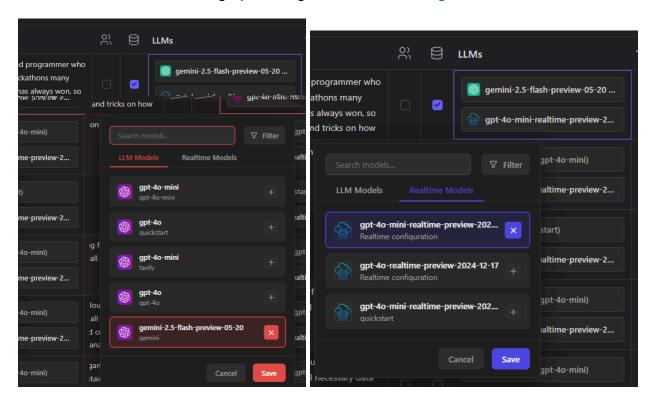


LLMs (Language Models)

This column shows which AI models are assigned to the agent. It may contain two subfields:

- LLM Models the main language model (e.g., GPT-4, Claude) used to process tasks for this agent.
- Realtime Models optional models that enable **real-time voice conversations** with the agent.

To find out more about setting up LLMs go here: <u>LLMs configurations</u>



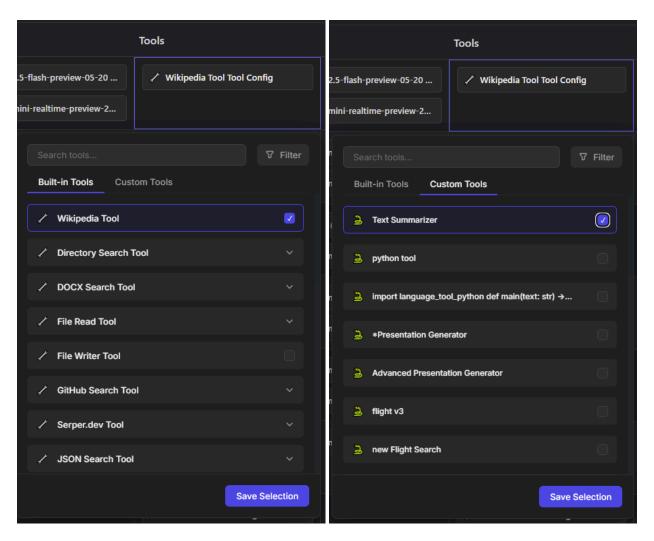
Tools

This column defines the capabilities the agent can use to perform its tasks.

- Built-in Tools tools that are part of the system by default (e.g., wikipedia, dall-e).
- Custom Tools Python-based tools created by the user. These can be added to give the agent access to custom logic, APIs, or special actions.

You can assign one or more tools to an agent using a dropdown or tool selector. Make sure the tools are relevant to the agent's role and goal.

Read how to create and manage tools here: <u>Tools</u>



At the end of each agent row, you'll find a **gear icon** . Clicking it opens a panel with **advanced configuration options** for the selected agent. These settings allow fine-tuning of how the agent behaves during execution.

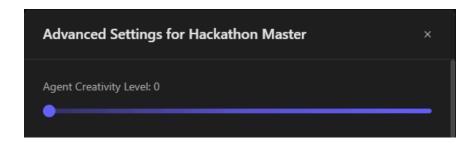


Here's what you'll find inside:

Agent Creativity Level (1–100)

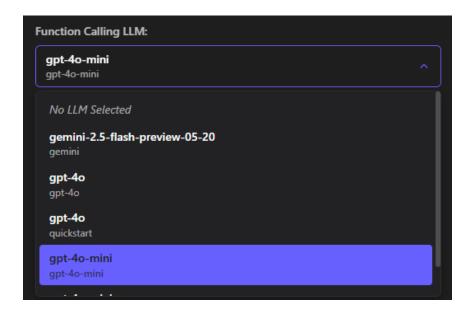
Controls how imaginative or focused the agent's responses are.

- → Lower values = more precise, conservative answers, does only what you ask of it.
- → Higher values = more creative, open-ended thinking.



Function Calling LLM

Select which language model should be used specifically for **function/tool calling**. Useful if you want to separate general reasoning from function execution.

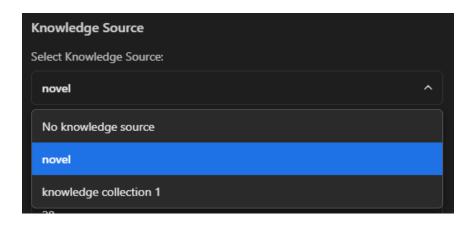


To find out more about setting up LLMs go here: <u>LLMs configurations</u>

Knowledge Source

Choose the data source the agent should rely on during execution — e.g., web search, internal docs, or no external knowledge.

Read how to create knowledge source here: Knowledge



Execution Settings

Fine-tune how the agent operates during a flow:

Maximum Iterations: Sets the maximum number of steps the agent can take to think, use tools, and generate results during one flow run.

Maximum Requests Per Minute: Throttles how often the agent can call tools or APIs.

Maximum Execution Time (seconds): Time limit for completing its tasks.

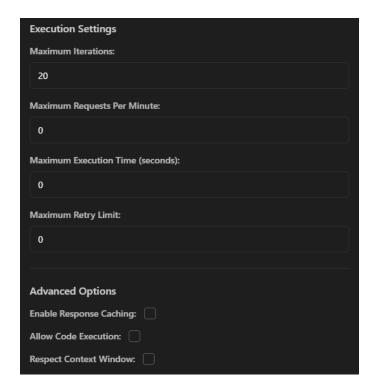
Maximum Retry Limit: How many times the agent will retry if something fails.

Advanced Options

Enable Response Caching: Caches the agent's responses to avoid repeating expensive operations.

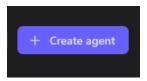
Allow Code Execution: Allows the agent to run Python or other executable code when needed.

Respect Context Window: Forces the agent to stay within the context limit of its model (prevents overly long responses or memory overflows).



2. Creating an Agent via the "Create Agent" Modal

Another way to create an agent is by clicking the "Create Agent" button in the top right corner of the page. This opens a **modal window** with a form for entering the agent's settings.

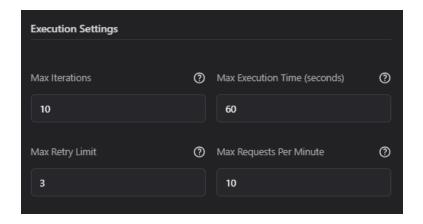


The form includes **almost the same fields** as in the table view — both the main properties and the advanced settings (normally hidden behind the gear icon).

Here are all the fields available in the modal:

- Agent Role
- Goal

- BackstoryAgent LLM



Notes on Creating Agents

To successfully create an agent, the following fields are **required**: **Agent Role**, **Goal**, and **Backstory**.

To ensure the agent can actually perform tasks in a flow, you also need to assign an **LLM model** (under **Agent LLM**).

All other fields are optional and can be customized later.

Table Actions (Right-Click Menu)

When working in the Table View, you can **right-click** on any row to access quick actions:

Add Empty Agent Above / Below

Inserts a new blank agent row directly above or below the selected one.

Copy Row

Copies the selected agent row (including all its settings).

Paste Row Above / Below

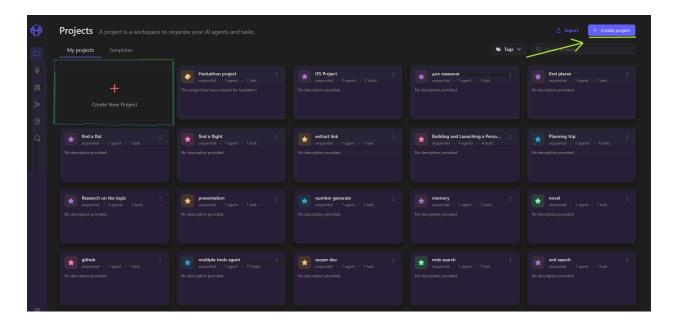
Pastes the previously copied agent into the table above or below the current one.

Delete Row

Removes the selected agent from the table.

Creating a Project

"Create Project"



modal window

Main Fields:

Project Name

Project Description

Process Type

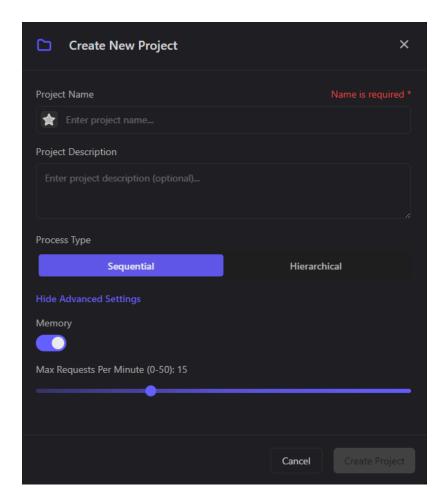
Sequential

Hierarchical

Advanced Settings:

Memory

Max Requests Per Minute



"Create"

Project Main Page

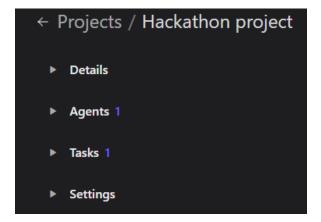
main page

Details

Agents

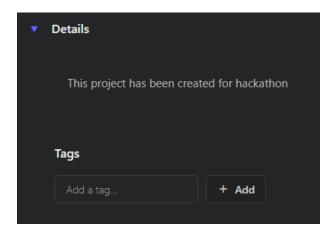
Tasks

Settings



1. Details

view or edit the project description



2. Agents

"Search agents"

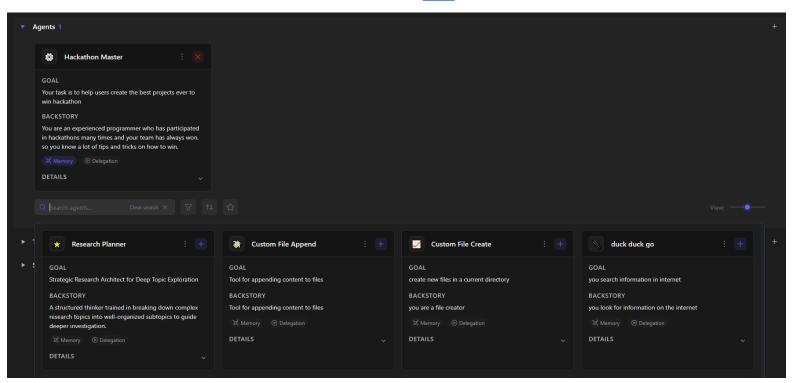


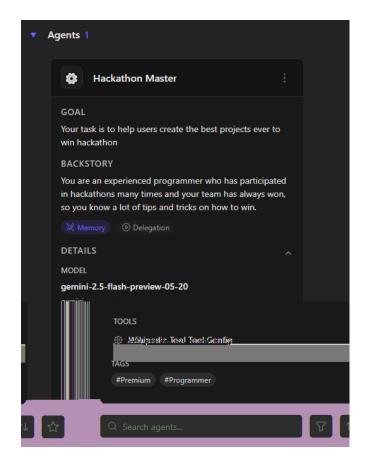
"Clear search"

"Details"



Staff





3. Tasks

Required fields to define a task:

Task Name Instructions

core prompt

task or action

Good instructions are:

Examples:

Expected Output

define the format or structure

Examples:



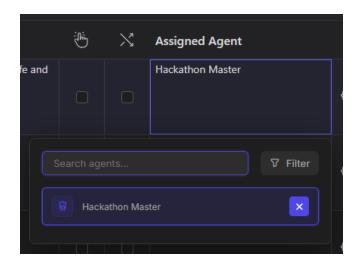
1 Important:

must be assigned to an agent cannot run without an agent

Assigning an agent to a task:

agent cell

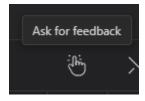
<u>Variables</u>



Additional options per task:

Human Input

Human Input



Advanced task settings:

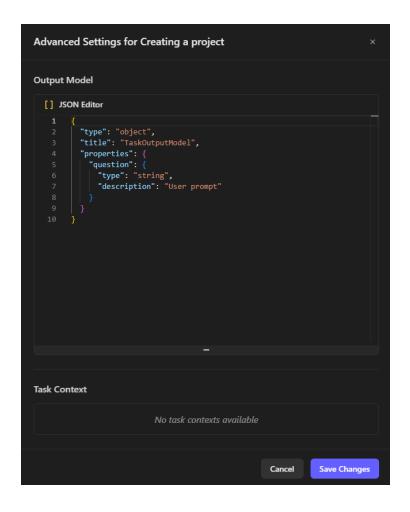
gear icon
Output Model

Advanced Settings

question type

string

description



Note

last task Output Model project node

4. Settings

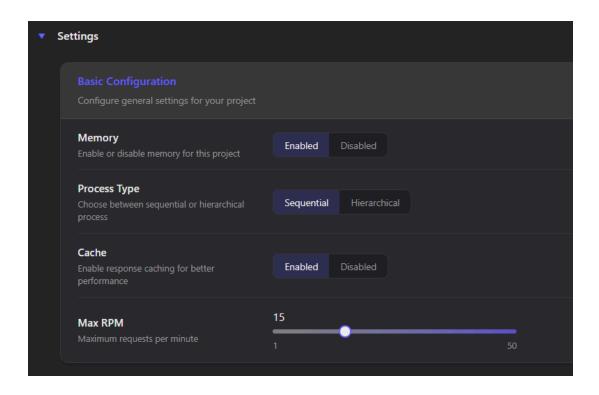
Memory

Memory

Process Type Sequential Hierarchical

Cache

Max Requests Per Minute (1-50)



What are tools — and why do agents need them?

Tools a a c a a ca ca d c d b d a a a d (LLM) ca d . T a plugins extra capabilities a d LLM.

LLM vs. Tool:

- T **LLM** (La a M d) a a:
 G a
 U d a d a a
 A a ba d a daa
 F c
- B LLM **cannot**:

 Ma API ca a da a

 D c a c d c

 R ad

 S a c b

 I ac da aba

Ta **tools** c .

Types of tools in the system

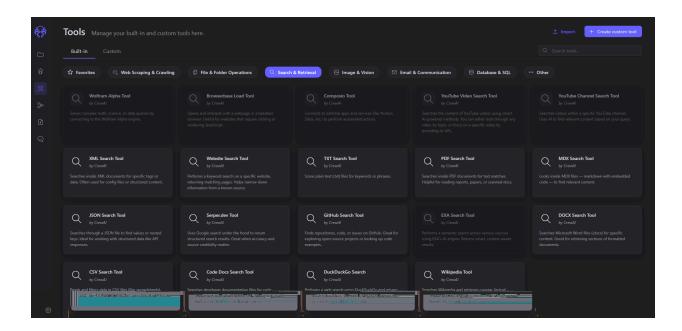
T a two kinds :

- 1. Built-in tools a ad a d(..,G Sac,Caca,Wb Sca)

Built-in Tools

B - a ad - ad a a ab .T a db ca (..,Sac,T P c ,C ca , c.)adcabaacda d ca ab .

H ca dadad aa : •••••

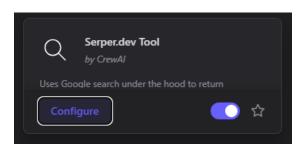


Tool Configurations

S add a c a .

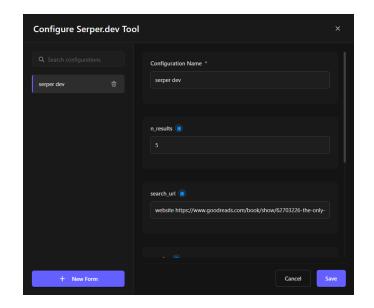
How to configure a built-in tool:

- 1. H .
- 2. C c "Configure" b a a a .
- 3. A c a



The configuration form includes:

- > A name for the configuration (d a)
- > Additional parameters (a b)



♠F d a blue info icon a ad a аа .



♠F d a d asterisks (*) a d ab Create

To add a new configuration:

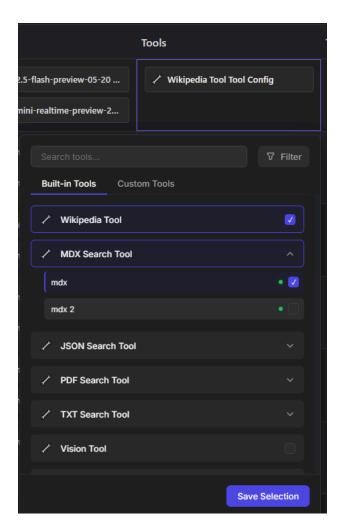
- 1. C c + **New Form**
- 2. F a a d d (a d *).
- 3. C c Create.

To edit an existing configuration:

- 1. S c c a .
- 2. Ma ca .
- 3. C c **Save**.

Assigning a built-in tool to an agent

- 1. G Agents c С.
- 2. D b -c c d **Tools** c d da .
- 3. I d , c **"Built-in"** ab. 4. C c a a .
- 5. I a a a a d c a , a dropdown menu a a.



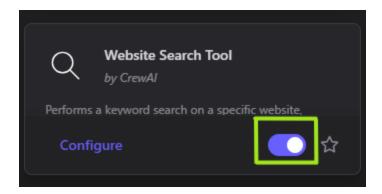
Disabling Tools

Y ca a **disable** a b c a a a

W d (ad), bc ac.

Disabled tools won't appear c d a

 $\mathsf{T} \qquad \qquad \mathsf{a} \qquad \mathsf{d} \quad \mathsf{ca} \; \; \mathsf{d}, \qquad \mathsf{d} \; \; \; \mathsf{d} \qquad .$



& Custom Tools (Python)

How to create a custom tool

1. G Tools a a d c c "Create Custom Tool" bc2. A da d d :

```
Name *

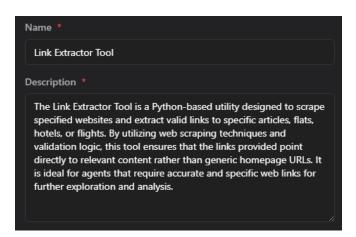
Link Estractor Tool

Description *

The link Estractor Tool is a Python based unity designed to scrape specific withouts and outside such districts the Indian State (as the Indian State). The outside such districts the Indian State (as the Indian State) is the Indian State (as the Indian State). The outside State (as the Indian State) is the Indian State (as the Indian State) is the Indian State (as the Indian State). The Indian State (as the Indian State) is the Indian State (as the Indian State) is the Indian State (as the Indian State). The Indian State (as the Indian State) is the Indian State (as the Indian State) is the Indian State (as the Indian State). The Indian State (as the Indian State) is the Indian State (as the
```

Required Fields:

Name а Description aa a d



Inputs Description:

```
Cc + c add a a
                      main() c
   Fac, d:
       Input Name
                                   P cd.
                   аc
                        a a
                              а
What it is: T Input Name
                                       Р
                   а
                       a aa
                              а
                                   С
    (main()) c . I must exactly match
                              a a
                                    а
                                             С
  а
Example: I c :
def main(city: str, date: str) -> dict:
Υ
                  a :
 • city
 date
       ac exactly ca - ,
```

Т

Description

What it is: A аа d а

Ea Dc :

- F city: "Destination city for flight search"
- F date: "Preferred departure date in YYYY-MM-DD format"

W ' : T d c a a (c \square), d a d a ac c

Ma a **Required**

Libraries:

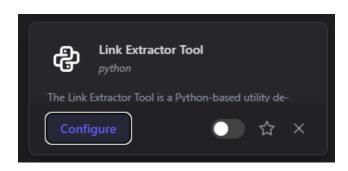
I cdd aba, "Enter library name..."

d

P Enter cc +b c

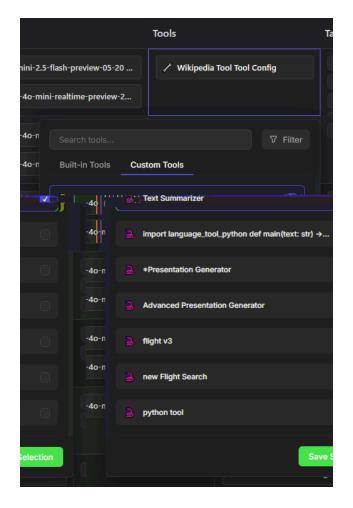
You don't need to add standard Python libraries manually — some are

A c a d a a Custom Tools ab T a T edit, a dcc Configure T delete, a d c c × (cross) c



Assigning a Custom Tool to an Agent

- 1. G Staff Table
- 2. D b -c c c d Tools c

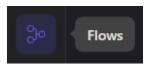


R ad c a a d c a : \underline{Staff}

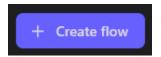
Flows let you automate how agents and tools interact in your project. Think of it as a visual pipeline where you define *what happens, in what order, and under what conditions*. Every project must be launched through a flow.

Step 1: Create a New Flow

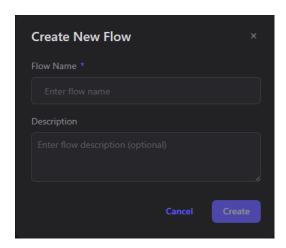
1. Go to the **Flows** page



2. Click the **+ Create Flow** button (top right corner)



- 3. A small window will open enter:
 - Flow Name (required)
 - Description (optional)



4. Click Create

You'll now be taken to the flow's canvas – a blank space where you'll build the logic of your flow.

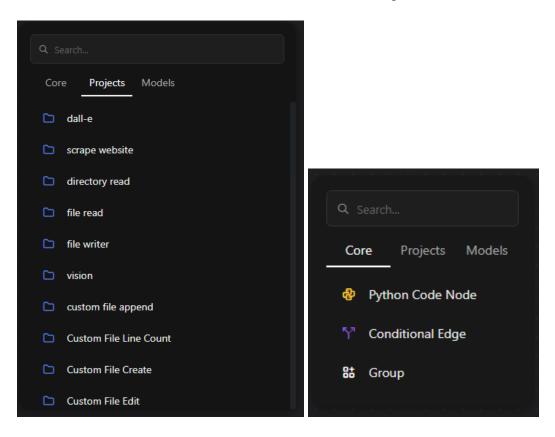
Step 2: Understand the flow canvas

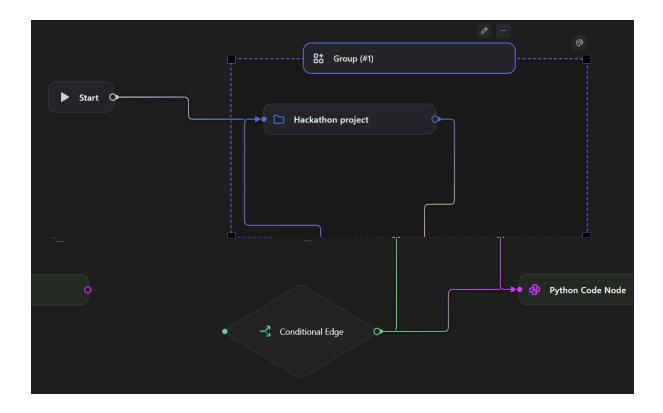
Each flow starts with a **Start Node** – it's added automatically and is always the entry point.

From there, you can add and connect different types of nodes:

- **Python Node** run Python code
- **Project Node** run previously created projects
- Conditional Edge set conditions (like if/else branches)
- **Group** organize your nodes visually

You can find out more about nodes and how to manage them here: Nodes





Step 3: Add nodes to the canvas

- 1. Right-click anywhere on the canvas
- 2. A popup will appear with node types
- 3. Navigate the tabs and click the type you want to add

The node will be placed on the canvas. You can **drag to reposition it** as needed.

For better navigation you can click on the magnifying glass icon, then you will see the list of all nodes present on your canvas. Click on any node and you will be redirected to it.

Step 4: Connect the nodes (Define the Flow)

- 1. Click and drag from one node's output port to another node's input
- 2. Arrows will appear between them
- 3. This defines the **execution order** from Start to Finish



Important: Your flow will run exactly in the order of these connections.

To delete a connection:

Click on the arrow and press **Delete** or use the **trash icon**



To delete a node: click it, then use the delete key or the trash icon.

Step 5: Configure each node

Each node has its own settings:

- Click on a node to open its configuration panel
- You can:
 - Set **variables** (e.g., pass data from one node to another)
 - Rename the node (each node must have a **unique name**)

We explain **variables** separately – but they let nodes exchange data and results. Here's the page dedicated to variables

Step 6: Save and run the Flow

Before running, you must save your flow:

1. Click **Save** (top right corner)



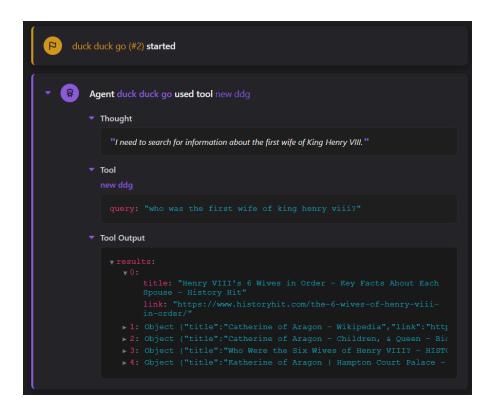
2. Then click Run

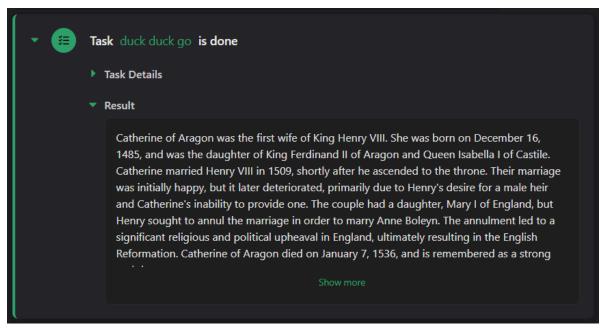
You will be redirected to the **Run Session Page**, where you can watch your flow execute live.

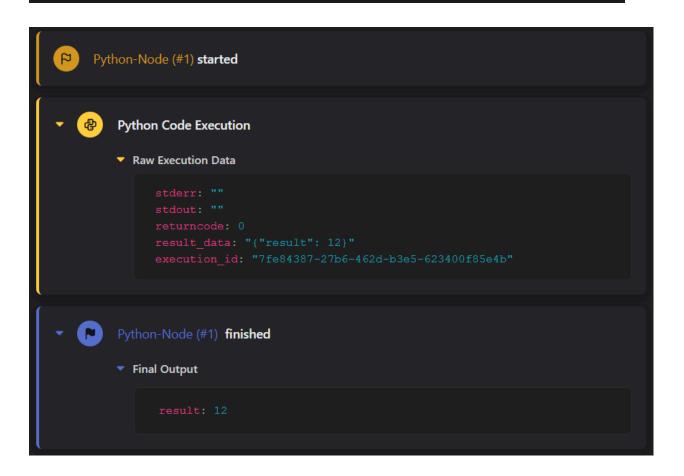
Step 7: What happens during execution?

As your flow runs:

- You'll see nodes activate one by one, in the defined order
- For each **Project Node** (agent task):
 - o You'll see:What the agent was asked to do
 - Their thought process (reasoning)
 - If a tool was used, you'll see:
 - The query
 - The tool output
 - Final result or decision







Execution Statuses – what do they mean?

At the top of the run screen, you'll see the **flow status**:



- Completed Flow finished successfully
- Running Flow is actively executing
- **Pending** Flow is waiting to start
- Waiting for Human Input Agent needs a user reply
- Error Something went wrong
- Expired Flow ran too long or was interrupted

Step 8: Reviewing sessions (past flow runs)

Each time you run a flow, it creates a **session** – a full log of what happened.

You can find past sessions in two places:

1. On the Run Page: click the **Sessions** button (top right)

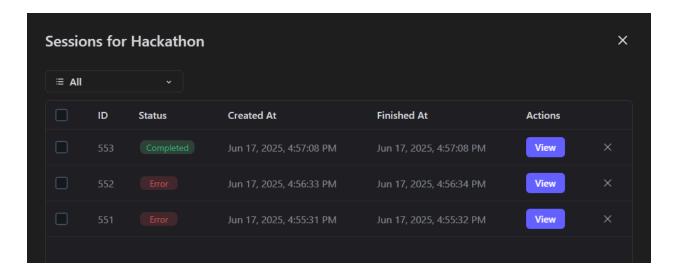


2. On the Flows list: click View Sessions next to any flow

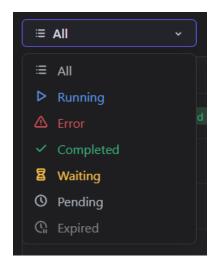


From the sessions panel, you can:

• Open any session to review it step-by-step



Filter by status



• **Delete** sessions using the X icon

Important Notes

- You can't run a project without a flow
- You must save before running otherwise your latest changes won't be applied
- Use clear and unique names for nodes to stay organized

Node Types in Flows

Each flow is made up of **nodes**, which are the building blocks of automation. Here's a breakdown of every node type, what it does, and how it can be configured.

1. Start Node

- This node is always present by default in every flow.
- It does not perform any action, but acts as the entry point of your flow.
- You can pass initial variables into it using JSON these variables will be available to all following nodes.

You can find a detailed explanation about variables here: <u>Variables</u>

2. Python Node

A Python Node lets you run custom Python code inside your flow. Use it when you need to:

- Transform or filter data
- Do calculations or logic
- Work with external APIs or data
- Prepare values for agents or later steps

Configuration Fields

When you add a Python Node, you'll see these fields:

Field	'What it does	
Name	The name of your Python tool (used in tool lists, not flow logic)	
Node name	A unique name for this node inside the flow (helps you identify it)	

Input map	Defines which flow variables are passed into your Python main() function
Output variable path	Where the result will be stored as a flow variable
Libraries	Python libraries your code needs. Press Enter/+ after typing each
You can find a -	
JSON	
Python	

This will multiply the values of variables.price and variables.quantity, and return the result.
To save the result, you must fill in Output variable path:
None
→ You can use variables.total_price in later nodes or agents.
Example 2 — Without input map:
If you don't use any variables or inputs, your code could look like this:
Python
In this case, the node always returns 42 — not dynamic or reusable.
Output variable path
This tells the system where to store your result from the Python code.
If your main() returns:
Python

And your Output variable path is:

None

Then variables.my result will contain the output and can be used by other nodes.

Summary

- Use Input map to pass values into main() from your flow.
- Your main() must return a dictionary like {"result": value}.
- Use Output variable path to save the result into a variable.
- If you don't use the input map, your code can still run but values will be hardcoded.
- Always press Save after making changes to the node!

3. Project Node

A **Project Node** lets you run any **existing project** inside a flow. It triggers a project you've already built in the system.

Configuration Fields

When you add a Project Node, you'll see the following fields:

Field	What it does
Node name	A unique name for this node inside your flow
Input map	Pass values from flow variables into the project
Output variable path	Save the final result of the project into a flow variable

How does it work?

When the flow reaches a **Project Node**, it:

- 1. Loads the selected **project**
- 2. If provided, uses the values from the **Input map** to populate variables inside the project
- 3. Executes the project step-by-step
- 4. Can save the final output into the variable you define in Output variable path

Example

Let's say your project expects a variable called city.

Input map:

```
JSON {
}
```

If variables.destination_city is "Paris", then the project will run using city = "Paris".

If the result of the project is a recommendation list, and you write:

```
None
```

in the **Output variable path**, then the output will be saved and can be used in later nodes.

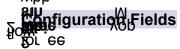
You can find a detailed explanation about variables here: Variables

4. Conditional Edge

ĮΧ

Fik

A **Conditional Edge** is a special type of node that helps your flow decide **which path to follow next**, based on logic written in Python.



whh

M

pas



How does it work?

1. In the **Python code box**, you write logic like:

```
Python

if variables.input_value > 10:
```

```
return "path_a"
else:
    return "path_b"
```

- 2. Based on what you return (like "path_a" or "path_b"), the flow will **follow the edge** with a matching label.
- 3. You can connect **two or more outputs** from the conditional edge to other nodes each with a **label** matching the values returned from the code.

How to set up the connections

- After writing your conditional logic, drag an edge from the output port to the next node.
- You'll be prompted to **enter a condition label** (like "path_a" "path_b" "default" etc.).
- The flow will follow the edge that **matches the returned label** from your Python code.

Example

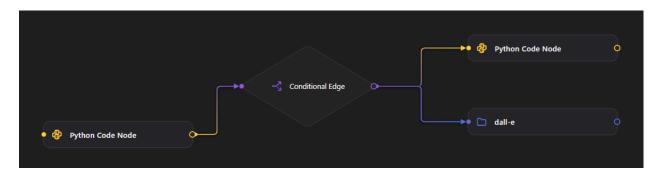
If your condition code is:

```
Python
if variables.status == "approved":
    return "go"
else:
    return "stop"
```

You should connect the conditional edge to two nodes:

- One with the label "go"
- One with the label "stop"

The flow will automatically route to the correct one based on the logic.



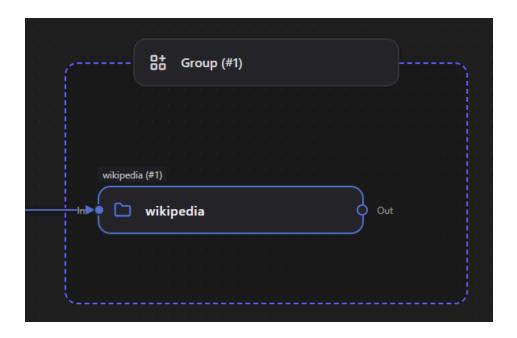
5. Group

A **Group** is a visual container inside your flow that helps organize nodes together. It's useful for keeping related logic in one place, especially in large flows.

How to use Groups

Add nodes to a Group:

- Just drag and drop any node into the group box.
- Once inside, the node becomes part of that group.



Remove a node from a Group:

- Hover over the node, and you'll see a special "remove from group" button
 - appear (usually near a corner).
- Click that button to take the node out of the group.

Group Options

You can customize the appearance and behavior of each group:

Feature	Description
Rename Group	Click the group name to edit it
Change Color	Use the color picker to change the background of the group
Collapse/Expand	Click the + / – button near the group name to hide or show its contents





Why use Groups?

- Keep your flow clean and easy to understand
- Visually separate different logic sections
- Make large flows easier to debug and manage

What are variables in flows?

Variables are a way to pass data between nodes in a flow.

They act like named containers that store values (such as text, numbers, or outputs from a node), and they make it possible for different parts of your flow to share and reuse data.

Where and how are variables used?

Variables are used in multiple parts of the flow system:

1. Start Node

- This is where the flow begins.
- You can set initial values for variables here.
- Example:
 In the Start Node configuration, you might define:

```
JSON
{
    "first_name": "Anna",
    "city": "Paris"
}
```

These variables can now be used by all other nodes.

More about nodes here: Nodes

2. Input Map

Every node (Python node, project node, conditional edge) has an Input Map field.

This is where you tell the node which variables to use as its inputs.

Example:

If your node expects a variable named first_name, your Input Map should look like this:

```
JSON
{
    "name": "first_name"
}
```

Inside your Python code or project, you can now access the variable as:

```
Python
variables.name
```

More about nodes here: Nodes

3. Output Variable Path

If your node returns a result, you can store that result in a new variable using this field.

Example:

• You write Python code that returns a value like:

```
Python
return {"result": greeting}
```

• In **Output Variable Path**, you type:

```
None
variables.my_result
```

Now, a new variable called my_result will be available for the next nodes.

More about nodes here: Nodes

Passing variables between nodes

Once a variable is defined or updated in one node, it can be passed along and used in the **Input Map** of any other node in the flow. This creates a **chain of data** between nodes.

More about nodes here: Nodes

Using variables in projects via {}

When working with **Project Nodes**, you can pass variables into the agent's task prompt or project template using **curly braces** {}.

Example:

If you defined a variable destination = "Rome" in the Start Node, you can write this in your project task:

"Find cheap flights to {destination}"

The system will replace {destination} with "Rome" automatically when the flow runs.

1 This only works if:

- The variable was already defined in a previous node (like the Start Node)
- It is available under variables.destination in the flow

More on how to create and launch project here: Projects, Flows

Variables in python code (main)

Inside a Python Node, variables are always accessed like this:

```
Python
input_value = variables.my_variable
```

You should never hard-code values inside your main() function unless it's for testing or temporary usage.

Correct way:

```
def main():
   name = variables.user_name
   return {"result": f"Hello, {name}!"}
```

Temporary (not recommended for final setup):

```
def main():
    name = "John" # Not using variables
    return {"result": f"Hello, {name}!"}
```

What happens if a variable is missing?

If you reference a variable that doesn't exist:

- The node might crash or return an error
- The flow could stop working at that point

Always make sure that:

- The variable is either defined in the **Start Node**
- Or it was created earlier in the flow using another node's **Output Variable Path**

Summary: Key Rules for Variables

Rule	Explanation
Variables must be passed using the variables. prefix	variables.city, variables.user_age, etc.
Input Map connects flow variables to node inputs	Tells the node what values to use
Output Variable Path stores a node's result	Makes it available to future nodes
In Project Node prompts, use {} to insert variables	<pre>Example: "Hello, {first_name}!"</pre>
In Python code, always access values through variables.	Not by hardcoding!
Variables are global inside the flow	Once created, they can be used by any future node

What is Human Input?

Human Input

pauses the flow

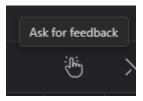
How to enable Human Input

per task

Here's how:

Project Page row

"Human Input"



This is the only way to activate Human Input.

Projects Flows

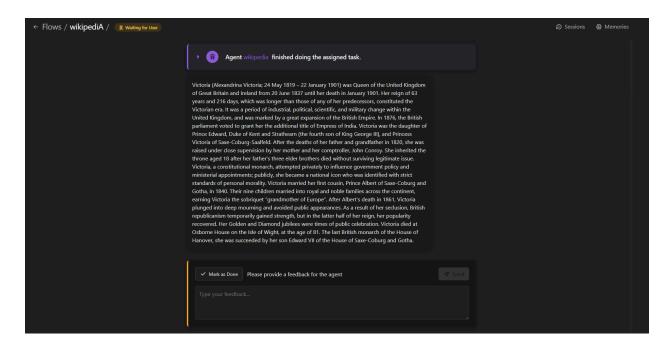
What happens during execution?

Human Input

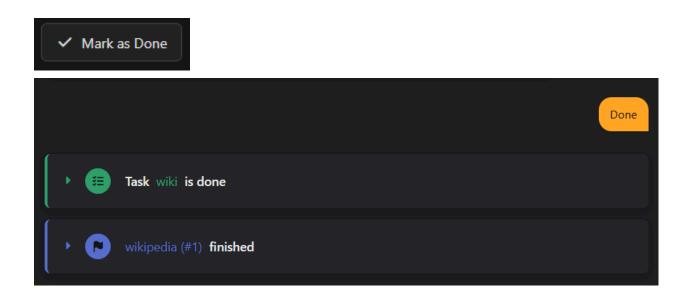
pause



chat window



type replies as many times as needed "Mark as Done"



1ECEV&•%ĐU Q€A•ð

togents

Star Constate ø

Agents can be enhanced with external documents and data through **Knowledge Sources**, allowing them to reference structured or unstructured information **before gaperiating responses**. This qowain sb ci i

0

Z L@H O@W V

Socializames i

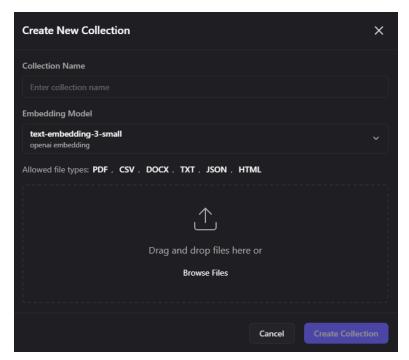
to conceee img

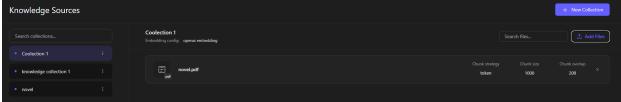
th

thingstiop.

plo

!I





File Settings: Chunking strategy

Once a file is uploaded, you'll configure how it's broken down ("chunked") for vector search.

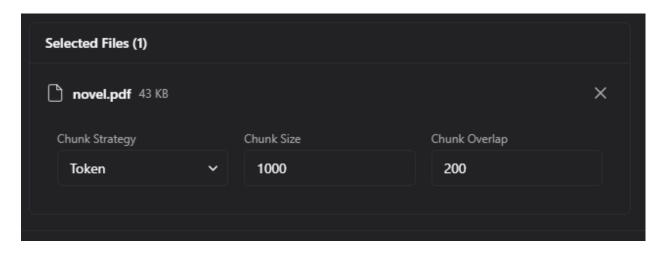
Parameters:

- 1. **Chunk Strategy** Choose how to split the content:
 - Token Splits based on number of tokens (language model units).
 - o Character Splits based on character count.
 - o Markdown Parses and splits using markdown structure (headers, lists).
 - o JSON Parses JSON objects.
 - HTML Parses structured HTML content.
- 2. **Chunk Size** The size of each chunk (in tokens, characters, etc., depending on the strategy). Affects how much content the model can reference at once.

Tip: Larger sizes = more context, but slower search.

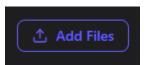
3. **Chunk Overlap** – Number of tokens/chars that repeat between chunks. Prevents splitting mid-sentence or mid-paragraph.





Managing files in a collection

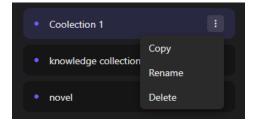
To add more files later:
 Click Add Files in the top-right corner of a collection.



To remove a file:

Click the **X** icon next to the file name.

To rename or delete a collection:
 Click the three dots beside the collection title.



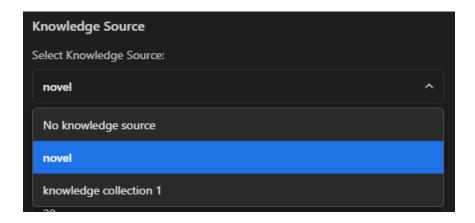
Assigning a knowledge source to an agent

Once your collection is ready:

- 1. Open the **Agents** table.
- 2. Click the **gear icon** (**) next to the agent you want to configure.



- 3. Scroll to the **Knowledge Source** section.
- 4. Under **Select Knowledge Source**, pick your desired collection from the dropdown.



5. Save changes.

Now the agent will reference that collection during execution.

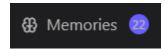
Read how to create agents here: Agents

Why Use Knowledge Sources?

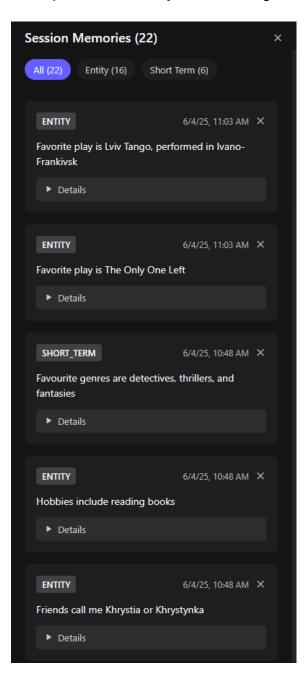
- Give agents access to internal data, guides, research, or instructions.
- Improve task accuracy and domain-specific reasoning.
- Avoid having to "teach" the same information in every prompt.

When your project is running in a flow (on the Run Session Page):

- 1. Look at the top-right corner of the screen
- 2. Click the "Memories" button



This opens a **sidebar panel** showing the memory contents for that session.



Types of memory

There are **four types of memory**, each serving a different purpose:

1. Short-Term Memory

- Automatically stores temporary thoughts or facts gathered by the agent during reasoning.
- Useful for immediate next steps.
- Example: intermediate steps in solving a problem.

2. Long-Term Memory

- Automatically stores important knowledge the agent decides is worth keeping during the session.
- Typically higher-level conclusions, goals, or facts.

3. Entity Memory

- Tracks structured information like:
 - Names
 - Dates
 - Locations
 - Other extracted data from text
- Example: if the user says "My name is Olivia", the system may store:

None

name: Olivia

4. User Memory

- Only created when a user **types something manually** using a Human Input node.
- Stores whatever the user enters, so it can be referenced later.

More about Human input here: Human Input

Interacting with Memory

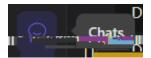
From the memory sidebar, users can:

- View each memory entry
- **Delete** any specific memory item
- **Filter** by memory type (or view all combined)

Summary

Feature	Description
Project Memory	Active, session-based, enables agents to remember info
Agent Memory	Not available yet
Activation	Enable in project settings
Access	Available via "Memory" button during a flow run
View/Edit	You can read or delete memory entries from the sidebar
Types	Short-Term, Long-Term, Entity, User Memory

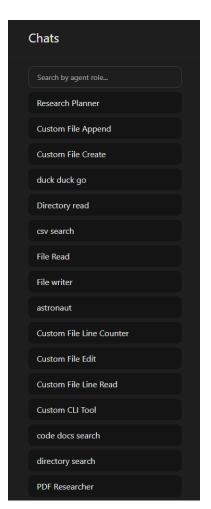
The **Realtime Chat** page allows you to talk to your agents **via voice or text** in real time. This feature is useful for testing agents interactively or having a natural conversation with them.



Agents Displayed

On this page, you will see all agents you've created in the system – the same ones listed in your **Staff Table**. (Read how to create and configure agents here: <u>Staff</u>)

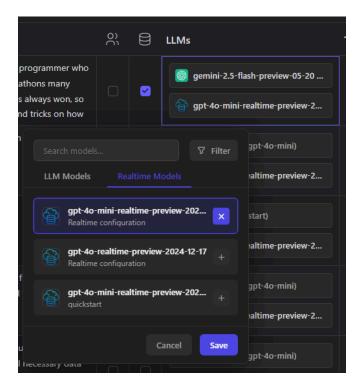
⚠ Only agents with a configured Realtime LLM model can respond in this chat.



Requirements to Enable Voice Chat

Before an agent can respond in Realtime:

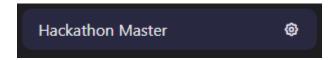
- 1. You must assign a Realtime LLM model
 - Go to the Staff Table
 - Double-click the LLMs cell of your agent
 - Open the Realtime Models tab in the popup
 - Click the + button to assign a Realtime model (only one can be active)

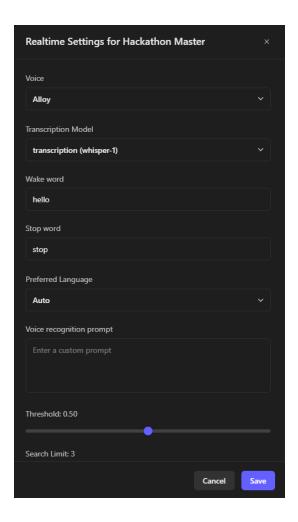


A detailed instructions on how to set up a realtime LLM is here: <u>LLMs</u>

Agent must be configured in the Realtime Chat page
 Next to the agent's name, click the gear icon to open their Realtime Settings.

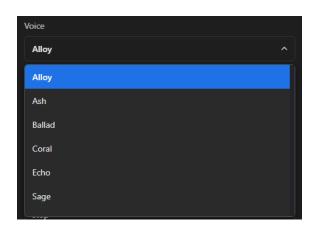
Realtime Agent Settings





In the Realtime Settings modal, you'll see:

• **Voice**: Choose a voice profile (cosmetic only – it doesn't affect functionality)



• Transcription Model ():
Used to convert your speech to text. Without it, the agent won't hear or respond.

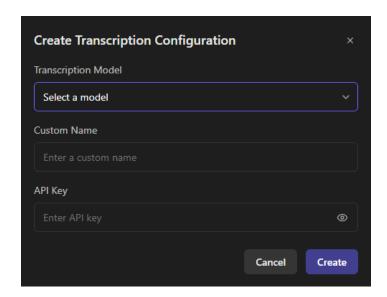


If no Transcription Model exists:

• Click the dropdown → Create New Configuration

Create New Configuration

- Fill out the modal form:
 - Transcription Model: Select one from the list (dropdown)
 - o Custom Name: Give it a name
 - o API Key: Add your service key if required



- Once the required fields are filled, the **Create** button becomes active.
- Click it to save and select your new configuration.

Additional Realtime Settings (Optional)

These options are shown in the agent's Realtime Settings window:

Field	Description
-------	-------------

Wake Word	A word that activates the agent (e.g., "Hey assistant", "Hello", Agent's name) Wake word hello	
Stop Word	A word that stops the agent from listening Stop word stop	
Preferred Language	Preferred Language Auto English Dutch Ukrainian French Arabic	
Voice Recognition Prompt	A hint to improve transcription accuracy (e.g., "You are listening to technical commands"). Also good practice to mention agent wake word here	
Threshold (0-1)	Threshold for searching by knowledge. Lower = more accurate knowledge will be extracted. Default: 0.6	
Search Limit (0-1000)	Limits how many knowledge chunks will be extracted. Default: 3	

How to Start Talking to Your Agent

Once the LLM and transcription settings are configured:

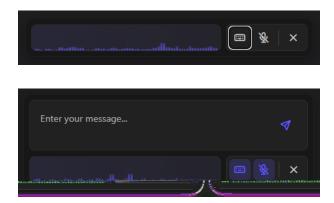
- 1. Allow microphone access in your browser when prompted
- 2. Select the agent you want to talk to
- 3. Click Start Speaking



Chat Interface Controls

During the chat session, you'll see:

Icon	Function
Microphone (when white)	Tap to speak – the system listens and responds
Microphone (when purple)	Click to mute yourself; the agent will not listen
Keyboard	Click to type a message instead of speaking
X (Close Chat)	Ends the conversation – all messages will be lost



Note: Realtime chats are **not saved**. Once the session is closed, all messages are gone.